Referee Report: Disemployment Effects of Unemployment Insurance: A Meta-Analysis - 2024-0441

Summary:

This paper conducts a meta-analysis of estimated effects of unemployment insurance. The authors collect elasticity estimates from the literature measuring the effects of either the replacement rate or potential benefit duration on unemployment/non-employment. These estimates tend to be closer to zero when the standard error is smaller, indicative of publication bias. Bias-corrected elasticity estimates are much closer to zero than a simple average of published estimates. The authors also show that estimates tend to be larger when the UI system provides greater benefits, either in terms of a higher replacement rate or longer potential duration, and that accounting for this element (as well as publication bias) substantially alters the optimal level of benefits in Baily-Chetty calculations.

This is an interesting paper on an important topic. Unemployment insurance has been extensively studied, which presents a valuable opportunity to synthesize results across studies. This paper uses well-established methodologies to conduct a meta-analysis and finds compelling evidence of publication bias. I think these results will be interesting to wide audience, not just that the corrected estimate is substantially different than the naive average, but also that this alters conclusions about optimal UI.

I've listed my comments below separated into major and minor comments:

Major Comments

- I found the exercise in Section 6 interesting, but I don't think it is exactly right to say that it is comparing the "macro" and "micro" effects. As the authors note, this exercise doesn't include any effect of UI on separations, for example through changes in aggregate demand. Also, there are three types of variation that are counted as capturing the "macro" effect; types (ii) and (iii) would not capture a macro effect since the control group consists of other workers in the same labor market. Type (i) could capture a macro effect since it's comparing to a prior cohort, but this then requires an identification assumption that the policy change is orthogonal to aggregate labor market developments. I think it would be more correct to say that the exercise tests the extent of spillovers or congestion in the labor market, which is an important channel through which the macro and micro effects could diverge, but not the only channel.

- If all of the studies in the meta-analysis were examining a single policy change (e.g. Card & Levine, 2000), then the heterogeneity analysis is Section 4 makes perfect sense to me. However, some studies are estimating an average effect across multiple policy changes (e.g. Farber et al. 2015 or Kroft & Notowidigdo 2016). This prompted two questions in my mind that would be helpful for the paper to clarify:
  - One, how are the economic characteristics assigned for studies that have multiple policy changes? I am guessing that you take the average across all of the settings studied, e.g. for Kroft & Notowidigdo take the average replacement rate across all US states over the years in their dataset. If this is the procedure, it would be helpful to state this somewhere.
  - Two, under what assumptions is it valid to include studies using multiple policy changes in the heterogeneity analysis? If the elasticity estimate you use for a study is implicitly an average across multiple policy changes, each of those changes could have different economic characteristics. In an ideal world, you could treat each of the different policy changes in that study as a separate estimate, and assign it the characteristics of that state/year. The usual logic says it should be possible to get the equivalent result as this policy-change-level regression by running a study-level regression where you've aggregated the individual policy changes using the same weights to aggregate the elasticities and the characteristics. So I think the implicit assumption being made is that the characteristics you're assigning to a study with multiple policy changes are representative of the variation that the study's regression places the most weight on. This could be violated, for example if the variation in Farber et al. (2015) came primarily from states with a higher-than-average baseline PBD, it would not be right to assign to this study the average US PBD. It would be helpful to somewhere state what assumption is needed to handle studies with multiple policy changes and explain why it is plausible in this setting.

Minor Comments

- In Figure 2, I might consider moving the right panel ("Andrews & Kasy 2019") to an appendix, or at least relabel it to make it more clear that this panel is not accounting for the elasticity heterogeneity estimated previously. Since you've shown that the elasticity increases with the replacement rate, I was initially confused why the right panel didn't show upward sloping lines, and I'm still not entirely sure what to take away from it.

- In Table 2, are there any additional institutional features of the labor market that predict heterogeneity in the elasticity? I'm thinking of measures of dynamism like labor market churn, average market thickness, etc., as well as measures of safety net generosity outside of UI.

- As an additional robustness check in Table B.4, have you considered estimating a specification in which $p(t)$ varies between three regions, $t < 0$, $0 \leq t < 1.96$, and $t \geq 1.96$? Anecdotally, I've heard skepticism of any even slightly negative estimates since those "go against theory". In Figure 1b, even the bottom tercile distribution is slightly right-skewed, which could be consistent with this.

- I typically think of publication bias as arising from researchers discarding insignificant results ex post, but in this setting could ex ante choices explain some of these findings? Researchers might be more focused on estimating elasticities in settings in which baseline UI benefits are relatively more generous, which you document tends to be associated with larger elasticity estimates.